

Big Data Summary



Map Reduce: Simplified Data Processing on Large Clusters¹ VS A Comparison of Approaches to Large-Scale Data Analysis²

By Cassandra Graves

¹Jeffrey Dean and Sanjay Ghemawat. 2008.
MapReduce: Simplified Data Processing on Large Clusters.
Commun. ACM 51, 1 (January 2008), 107-113.
DOI=10.1145/1327452.1327492
<http://doi.acm.org/10.1145/1327452.1327492>

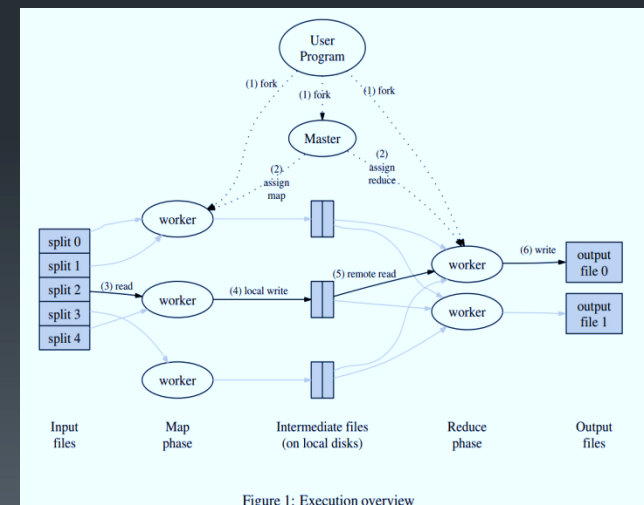
² Andrew Pavlo, Erik Paulson, Alexander Rasin, Daniel J. Abadi, David J. DeWitt, Samuel Madden, and Michael Stonebraker. 2009.
A comparison of approaches to large-scale data analysis.
In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data (SIGMOD '09)*,
Carsten Binnig and Benoit Dageville (Eds.). ACM, New York,
NY, USA, 165-178. DOI=10.1145/1559845.1559865
<http://doi.acm.org/10.1145/1559845.1559865>

Main Idea About MapReduce

- MapReduce is a programming model that is easy to use
 - Programmers without experience with parallel and distributed systems can easily use MapReduce
 - Hides parallelization, fault-tolerance, locality optimization and load balancing from the programmer
- Users must specify two functions
 - *Map*, function written by user, takes an input pair and produces a set of intermediate key/value pairs. All intermediate values associated with the same intermediate key I are grouped together
 - *Reduce*, function written by user, accepts an intermediate key I and a set of values for that key. Values are passed via iterator, which allows us to handle lists of values that are too large to fit in memory
- MapReduce automatically parallelizes and executes on a large cluster of commodity machines
- MapReduce uses redundant execution to reduce the impact of slow machines as well as to handle machine failures and data loss

Implementing MapReduce

- Implementation varies - based on environment being used upon
- Example implementation:
 - Dual-processor x86 machines running Linux, with 2-4GB of memory per machine
 - Commodity networking hardware
 - Cluster has hundreds or thousands of machines
 - Machine failures expected
 - Storage using IDE disks on each machine and a distributed file system to manage data
 - Submitted jobs pass through a scheduling system
 - Each job's set of tasks is assigned to available machines within the cluster
- 7 Step Execution
 - The MapReduce library splits the input files and starts copies of the program on a cluster of machines
 - One copy becomes the master and it assigns work to the remaining
 - A worker reads the input, parses key/value pairs to pass into the Map function and are buffered in memory
 - The buffered pairs are written to local disk, partitioned into R regions and forwarded to the reduce workers
 - Reduce workers read the buffered data from the local disks, sorts the data by key to get all occurrences of the same key together to do same reduce tasks
 - Reduce worker passes the key and corresponding values to the Reduce function and the output gets stored in a final output file for this reduce partition
 - Once all map and reduce tasks are done, control is given back to the user code



Analysis

- MapReduce appears to be an easy tool to use when looking to distribute work over a cluster of machines
- Since it automatically takes care of coordination between machines in the cluster, that would improve productivity of the programmer
 - More time to add to the tasks versus taking up time coordinating
- Graceful fault-tolerance in MapReduce is a motivating feature, as the program won't crash if one machine fails
 - As this takes time, it would be wasteful to have to restart the entire program if one machine were to break down
- Use of local storage on machines after *Map* seems helpful, but if a machine only finished *Map* when it fails, then *Map* must be done again
 - Seems redundant and counterproductive if/when many machines fail
- Redundant execution is used towards the end to ensure completion time isn't halted by a slow machine
 - Seems like a good idea, but a lot of overhead is being added into the program when in most cases, it might not need to
- Google uses it – so it must be good.

MapReduce vs Comparison Paper

- Comparison paper supports claim of MapReduce's simple model to utilize a cluster of machines
 - Mentions specifically as a learning tool
- Parallel DBMS also provide a high-level programming environment
- Parallel DBMS need a well-defined schema for data, instead of a schema, MR needs a custom parser to gather data
- Unlike Parallel DBMS' use of hash or B-tree indexes, MR does not provide indexing, which is complicated to implement and maintain
- DBMS use straight forward requests for information versus MR needing to provide an algorithm to request data
- Basic database tasks take significantly longer to load and execute for MR then it does for Parallel DBMS
- MR performs complete table scans, while DBMS takes advantage of clustered indexes – MR deals with more overhead
- MR has a slow-to-begin “cold start” versus the quick-and-ready “warm start” that DBMS has
- Both are easy to use, but MR is difficult to maintain once beyond beginner concepts

Advantages/Disadvantages

■ Advantages

- With independence of each machine in the cluster and no schema required, MapReduce is very flexible
- Creation of high-level languages Pig and Hive to alleviate implementation of repetitive tasks done with MapReduce
- Best implementation of fault tolerance
- Less challenging to install and configure properly MapReduce(Hadoop)
- Minimizes work lost when a machine in the cluster fails

■ Disadvantages

- With each machine being independent in the cluster, and not having a defined schema, MapReduce could easily be corrupt by bad data
- Forced to write algorithms in a low-level language in order to perform record-level manipulation in MapReduce
- Reduce creates multiple output files, so another instance of Reduce would be necessary to combine them all into one file
- Load and execution time of basic tasks take significantly longer time for MapReduce(Hadoop) versus DBMS (Vertica & DBMS-X)
- Extensive amount of overhead